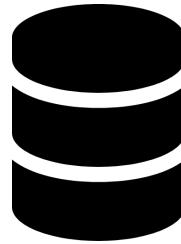


Microservices

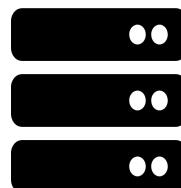
(Slides prepared by Derek Brown)

1. A Short History of Applications at Scale

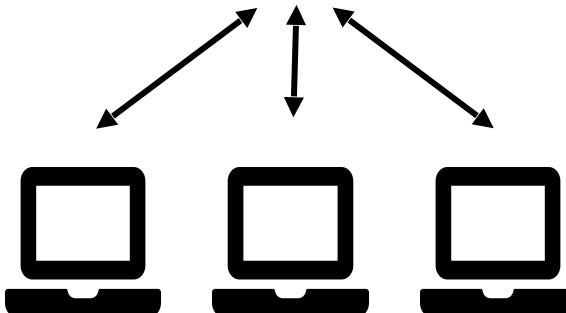
Motivation: The Early web



Database Process
(SQL)



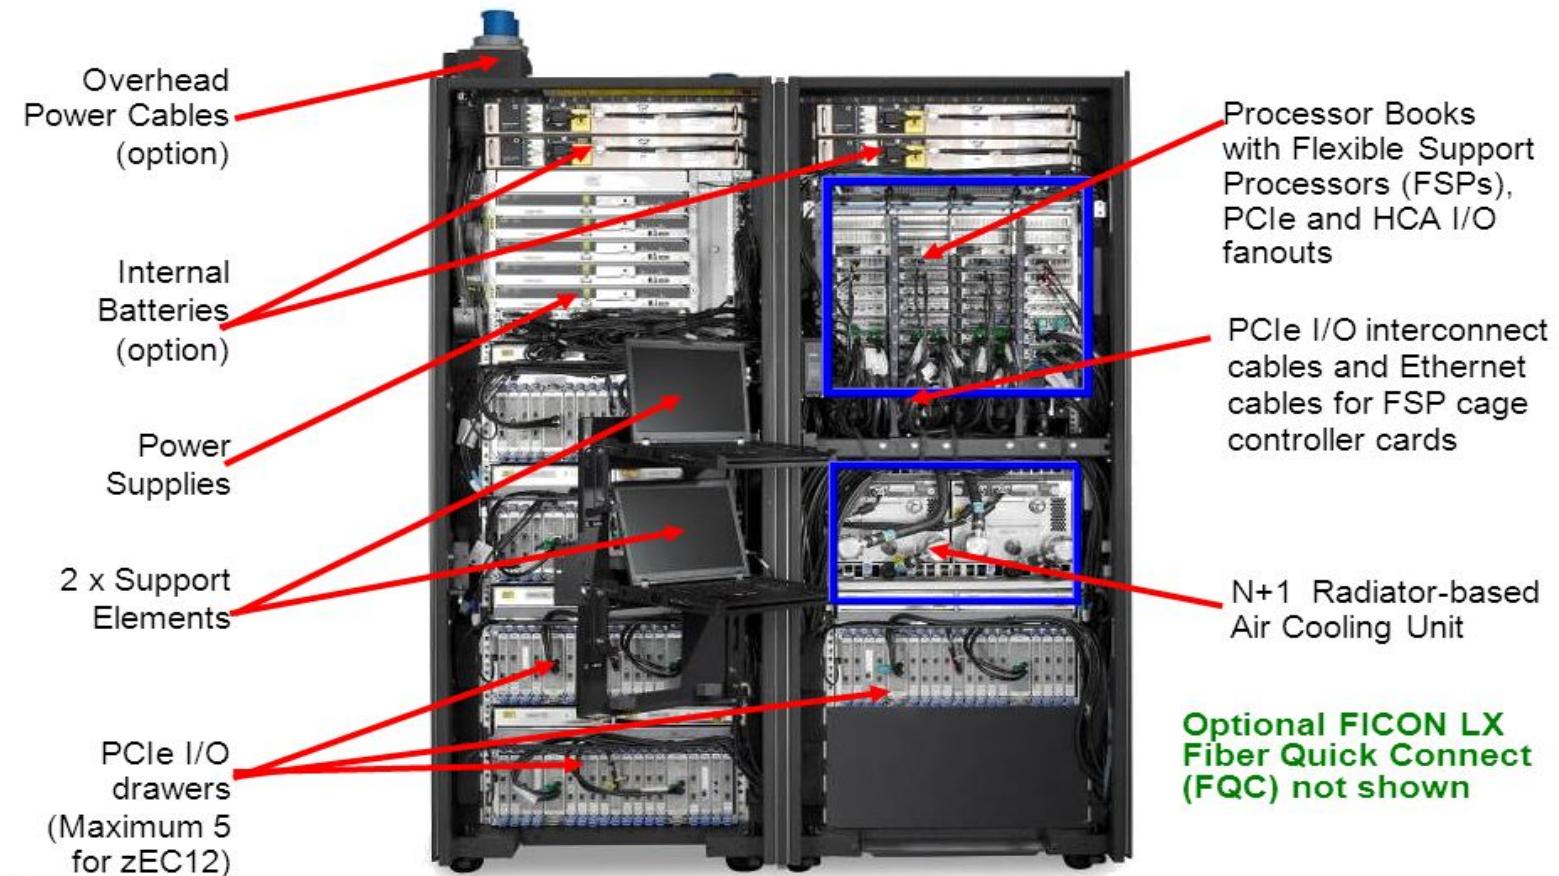
Single, Monolithic App Server
(PHP, Rails, Java EE)



Limited Clients, Mostly Text Data

Q: All the sudden, we have a lot more clients. What do we do?

Q: All the sudden, we have a lot more clients. What do we do?



MAINFRAME!

Motivation: The Early Web

Advantages of the “Mainframe Model”:

- Simple deployment model
- Easy to manage
 - Single point of entry for security / network access.

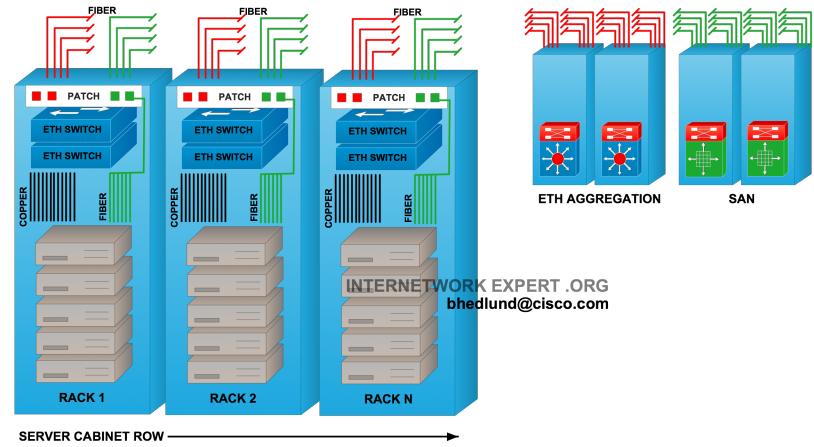
Motivation: The Early Web

Disadvantages of the “Mainframe Model”:

- Cost increases Exponentially with Scale
 - Increased power density
 - Requires smaller circuitry / higher operating frequencies
 - Complex
- Low Redundancy
- Resource Sharing Inefficiencies

Q: Eventually, single servers couldn't be made any faster. So how do we continue to serve more clients?

Q: Eventually, single servers couldn't be made any faster. So how do we continue to serve more clients?



DATA CENTER!

Aside: Not all servers are created equal



Block / Object Storage Servers

Store data for other applications like Databases.



General Compute

Good all-around servers

GPU Servers

Perform parallel processing
(ex. Machine Learning)



Appliances / Middleboxes

Servers with preloaded software to serve a dedicated purpose.
(ex. Firewall, SDN, VM Server)

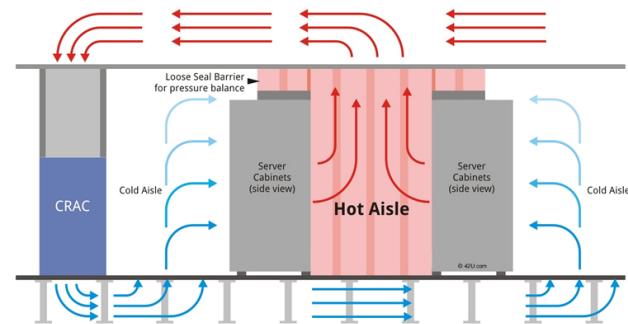
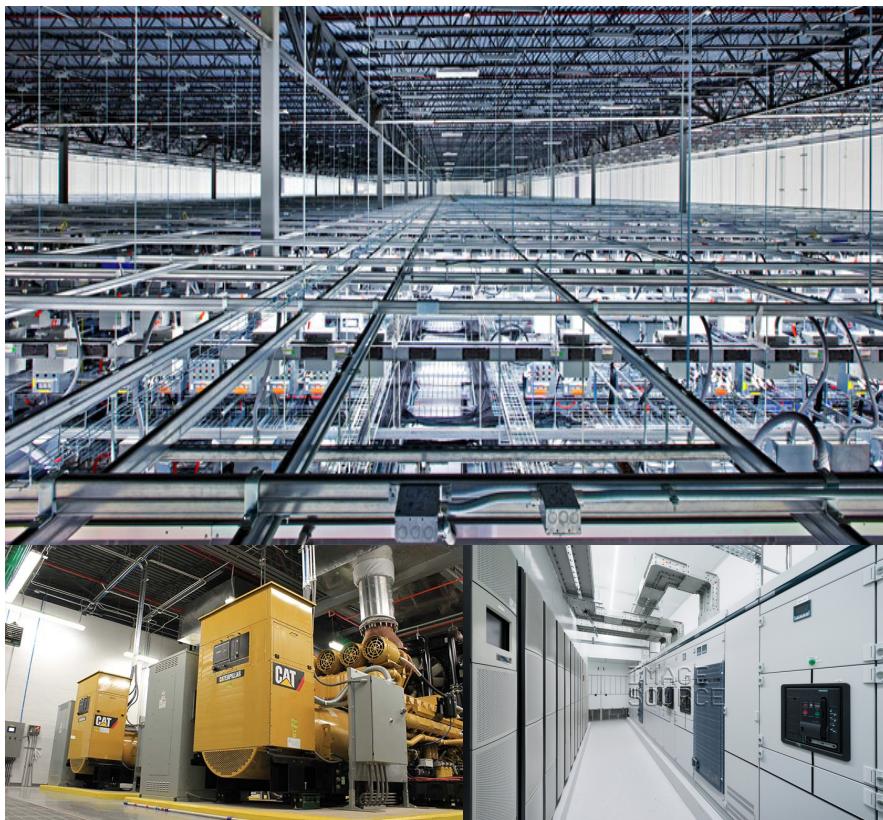
Blade Servers

Small dedicated instances for clients / applications.

Remote Access Servers

Small servers embedded in big servers for provisioning and control.

Motivation: Datacenters are complicated



- Complex network setup (Single/Multi Fiber, Network Switches, External Switching)
- Computer Room Air Cooling system (CRAC) to maintain operating temperature.
- Uninterruptible Power Supply (UPS); multiple power “grids”, multiple power “phases”, battery backups for faults, diesel generators.

Motivation: Datacenters

Advantages of the “Datacenter Model”:

- Scaling is relatively easy; just buy more servers.
- Fairly Redundant
 - Multiple servers can perform the same task.
 - Disks and Power supplies are “hot swappable”, allowing servers to run even while maintenance is being performed.
- Application residency; applications can be put on a specific set of servers based on resource requirements.

Motivation: Datacenters

Disadvantages of the “Datacenter Model”:

- Applications must be distributed/scalable
 - Hard to distribute data across multiple instances.
- Management is difficult
 - Network topologies can be complex
 - Automatic installation of OSes/App is hard
- Must physically provision servers

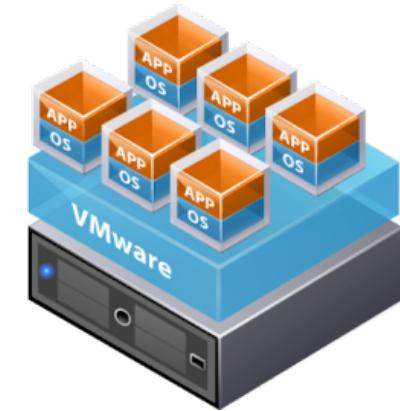
And... Servers can be Expensive

Item	Quantity	Price
	<input type="button" value="1"/> ▼	\$393,350.00
<p>New PowerEdge R940 Rack Server</p> <p>Edit Save for later Remove</p>  <p>BUYS \$247K SERVER</p> <p>USES IT TO PLAY MINECRAFT</p>		<p>-\$145,827.04</p> <p>\$247,522.96</p>

Q: WHAT IF I DON'T WANT A SEPARATE SERVER FOR EACH APP?



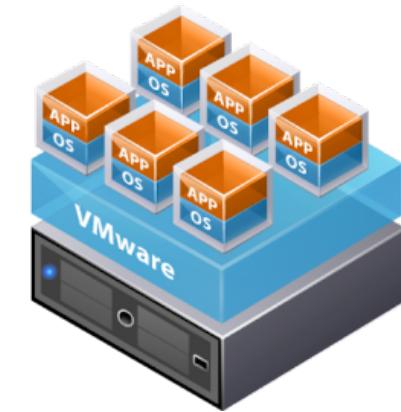
Microsoft
Hyper-V



Q: WHAT IF I DON'T WANT A SEPARATE SERVER FOR EACH APP?



Microsoft
Hyper-V



VIRTUAL MACHINES!

Motivation: Virtual Machines

Advantages of Virtual Machines

- Can divide servers into separate smaller pieces
 - We don't need to devote a ton of resources to each app.
 - Means we can “over-schedule” servers.
 - Applications can run more securely as they are “sandboxed”
 - Allows for more complex virtual networking.
- Allows for hardware portability.

Motivation: Virtual Machines

Disadvantages of the Virtual Machines:

- Adds additional layers of abstraction
 - Reduces performance of system as emulation is slow.
 - Adds more complexity in terms of scheduling.
- More expensive (VMware ESXi ~\$5k/server)

Q: I am a developer. I don't want to have to think about what physical server is hosting my app. How do I get just a generic server instance?



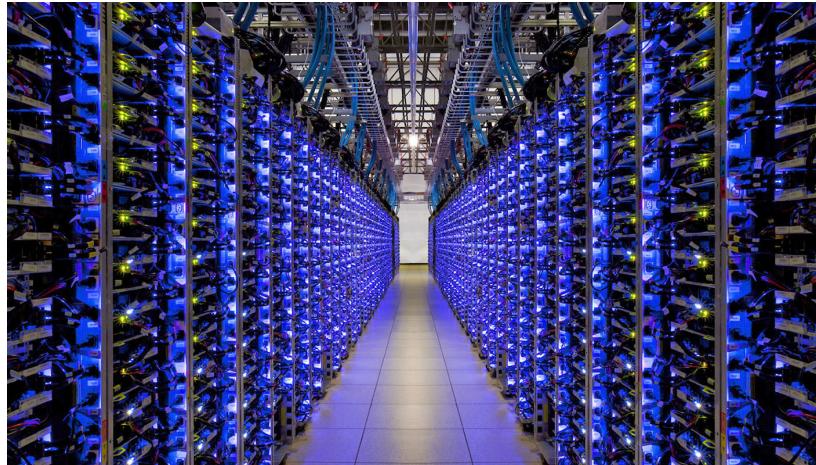
Google Cloud Platform

vmware®

openstack.

Azure

amazon
web services



Q: I am a developer. I don't want to have to think about what physical server is hosting my app. How do I get just a generic server instance?

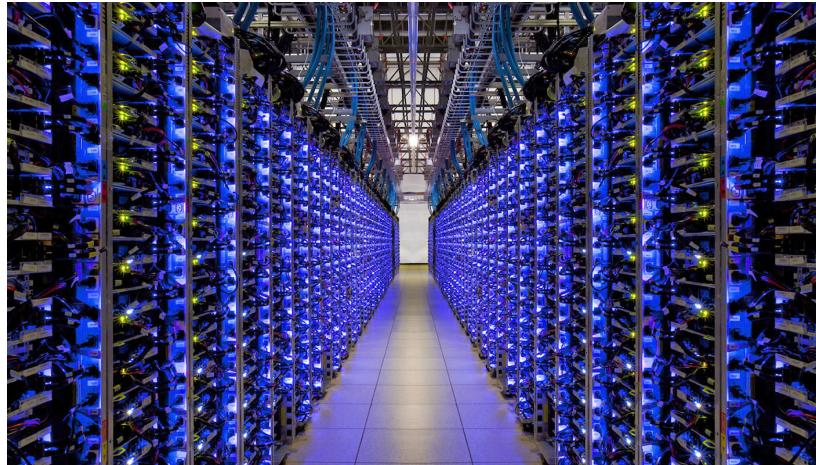


Google Cloud Platform

vmware®

openstack.

Azure



**amazon
web services**

CLOUD!

Aside: A Tale of Two Clouds



Google Cloud Platform



Public Cloud

- Hardware managed by a 3rd Party
- Product is the server as a service

Private Cloud

- Hardware is managed internally.
- Product is the management platform

Motivation: Cloud

Advantages of the “Cloud”:

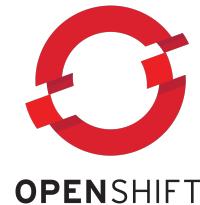
- Really easy to provision an instance (as long as resources exist)
- Don't have to think about hardware
 - Can move “virtual machines” across systems
 - Can setup complex virtual networks
 - Can divide a server into smaller virtual instances
 - Automatic backup and replication
- Can over-schedule servers

Motivation: Cloud

Disadvantages of the “Cloud”:

- Requires substantial overhead for VMs/management layer.
- Less control over data residency
- Other applications/customers can affect your application

Q: I am a developer. I don't want to have to think about servers at all. I just want to write code!

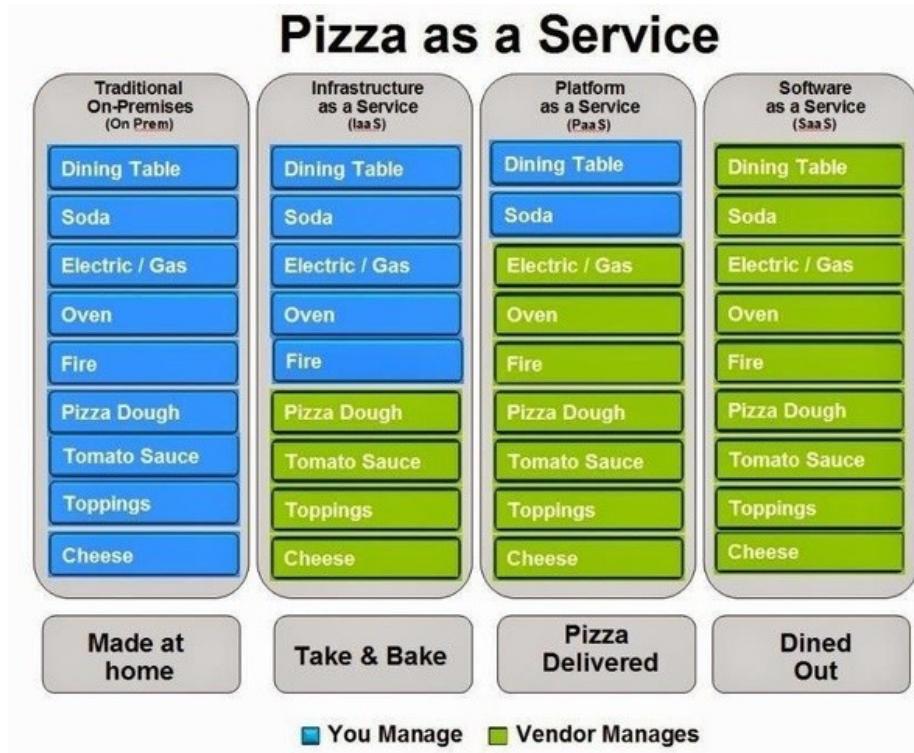


Q: I am a developer. I don't want to have to think about servers at all. I just want to write code!



PLATFORM AS A SERVICE!

Recap: By Analogy to Pizza



We are implicitly trading control for ease-of-use.

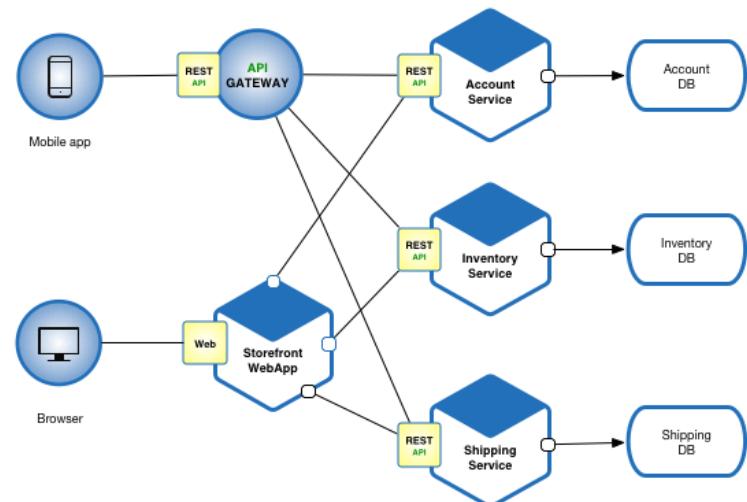
2. What are Microservices?

KEY IDEA:

The Microservice architecture is a method of developing distributed applications.

Microservice Philosophy

- Microservices are applications organized around a single capability.
- Small in size, modular, replaceable, independently deployable.
- Stateless- meaning they are scalable.



Aside - Stateless Programming

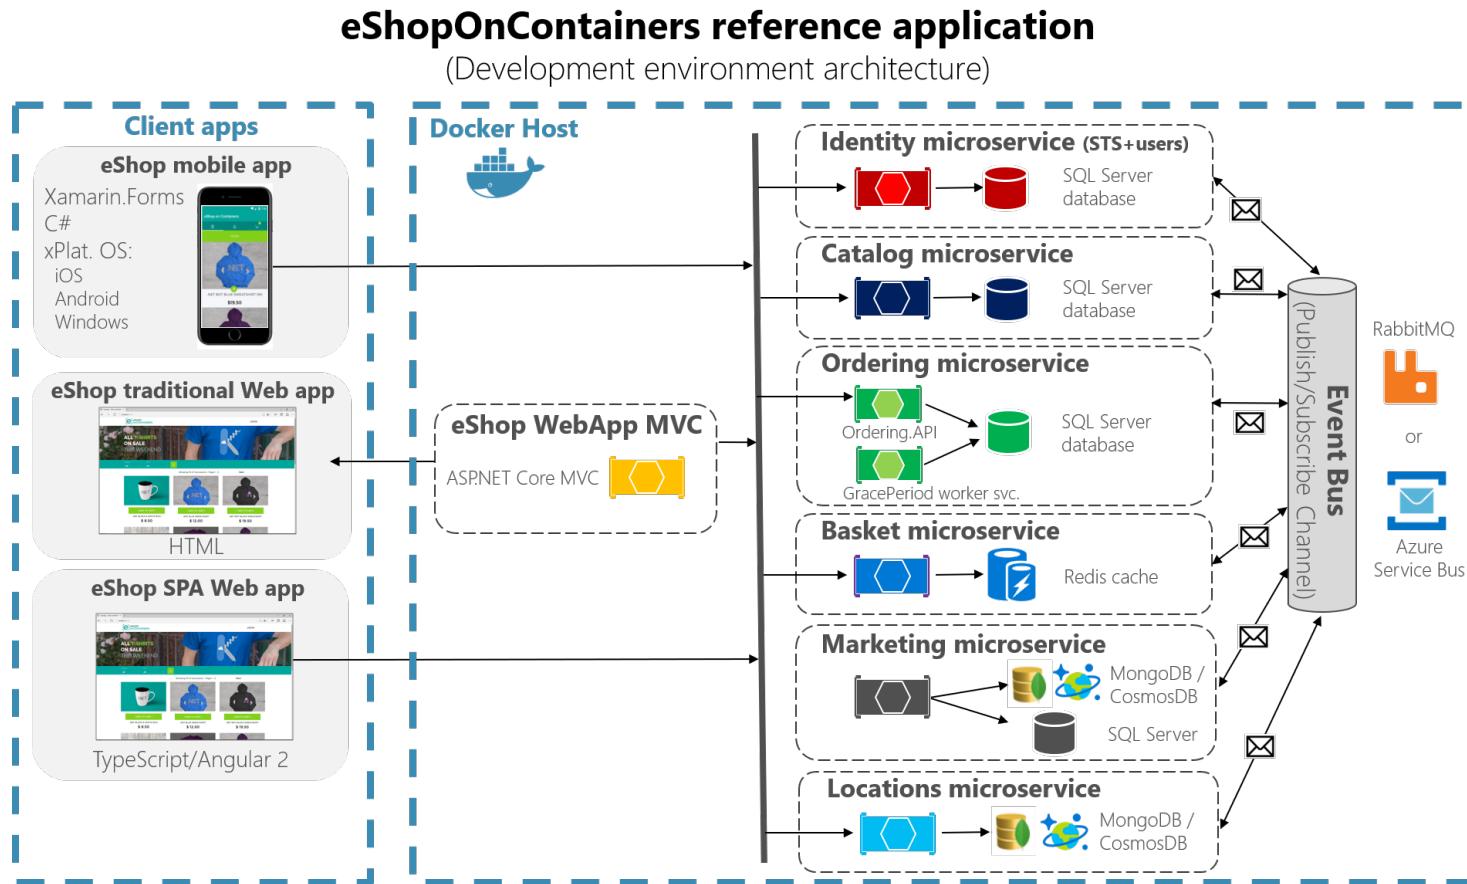
Stateless means there is no record of previous interactions and each interaction request has to be handled based entirely on information that comes with it.

SOURCE: <http://whatis.techtarget.com/definition/stateless>

Why?

- Microservices are continually created and destroyed. Data stored in an instance can be erased.
- Microservices are built to be scaled (in parallel). Storing stateful data can lead to race conditions.
- This offloads computation from the server to the client, speeding up the application.

Example: Mobile Shop



SOURCE:<https://docs.microsoft.com/en-us/dotnet/standard/microservices-architecture/multi-container-microservice-net-applications/microservice-application-design>

Class Exercise - Netflix

What types of microservices would you expect Netflix to run in production?



Class Exercise - Netflix

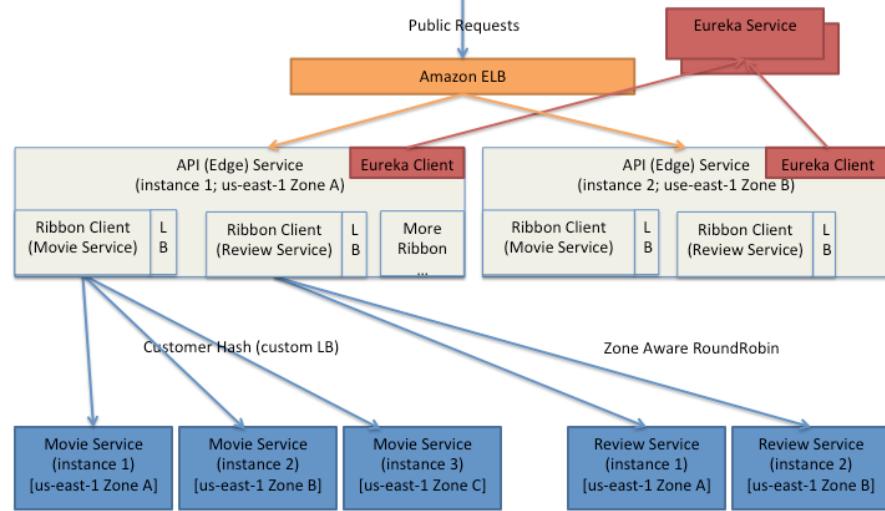
What types of microservices would you expect Netflix to run in production?



Class Exercise - Netflix

Services Netflix has:

- API Gateway Service
- Titles Service (for Metadata)
- Ratings Service
- Movies Service (for Data)
- Accounts Service
- A/B Testing Service
- Service Failure Service
- Rescheduling Service
- Log Collection Service
- Service Status Visualization Service
- Transaction Service
- ...



3. How does the Microservices Architecture affect software development practices?

Class Exercise

How would you expect microservices would affect:

1. Performance?
2. Error Tolerance / Recoverability?
3. Ability to find a bug in the codebase?
4. Team Size?
5. Length of development/release cycle?
6. Willingness to adopt new technologies?

Class Exercise

1. Performance

- Likely to increase as the application scales more flexibly.
- However, poor architecture (such as a bottleneck) reduces efficacy.

2. Error Tolerance / Recoverability

- Can be less fault tolerant, as a single failed service can cause entire application to fail.
- But this can be mitigated by redundant services and fallbacks.

3. Ability to find a bug in the codebase

- Much easier. Errors are segmented by service, reported directly to team.

4. Team Size

- Likely smaller, as less integration is required.

5. Length of development/release cycle

- Services are likely quick to develop, as they can be created on a blank canvas.
- Once created it can be difficult to refactor all the API, making change slow.

6. Willingness to adopt new technologies

- Likely more willing to adopt new technologies, as each service can operate on its own tech stack.